M3 Maths Challenge: Staying Cool as the World Heats Up

Team #17559

1 March 2025

1 Executive Summary

Indeed it is no secret that global warming is going to have massive implications on global temperatures, a cities' economy and the overall health of a city's population. In this paper we explore these 3 domains and analyse the extent of the effect that global warming has using various mathematical models implemented in the Python programming language.

Our team investigates how global warming is affecting Birmingham, England, currently and years to come using the predictions from our models and by analysing past data.

In Q1 we were asked to predict the indoor temperatures for resedential homes in without air conditioning over a time period of 24 hours during a heatwave. We use the dataset given [2] to create a multivariate polynomial regression model of degree 3 in order to predict indoor temperatures t seconds after midnight. We have developed a Python program that implements a function called `Predict_Temps(new_data, year)` that predicts temperatures over the course of 24 hours and produces a list of predicted temperatures across each hour in the entire 24 hours. A complete concrete list of predicted temperatures per hour can be viewed at the end of section 2.5.

In Q2 we were asked to predict the peak demand of energy consumption in kWh now and 20 years from now. We implement the Holt-Winters exponential smoothing model in Python in order to make predictions for these values. From our predictions we conclude that the peak energy demand that Birmingham should be prepared to handle now is 2.76×10^8 kWh, and in 20 years from now the value will be 2.68×10^8 kWh.

Finally in Q3 we were asked to find a suitable 'vulnerability scale' that could be assigned to any sub-city within Birmingham which would then allow city officials to determine the amount of resources allocated to these particular areas in the event of a heat wave or power grid failure. We decided to calculate the heat index values via Python in order to model how these values would look for the actual heat wave for the data that M3 provides. From this data, we can suggest to city officials that some of the power of the UK, given that they are affected to a lesser extent by the heat waves, should be directed to irmingham during the heat wave in order to allow most of the population, i.e. the ones who have access to AC, to activate it to minimize the effects of any heat stroke and also by providing this excess power it prevents any power outages or power grid failures.

Contents

2 Q1: Hot to Go

2.1 Problem Statement

In Q1 we interpret the problem statement like so:

Predict indoor temperatures for resedential homes without air conditioning over a time period of 24 hours, in Birmingham England, during a heatwave.

2.2 Assumptions

We assume that all homes are made from bricks or another similar material.

Justification: For this question we have selected to analyse homes in Birmingham, England. It is intuitive to reason that all homes are built out of bricks, concrete or some other similar material. Indeed we understand that the United Kingdom is considered to be an advanced developed country, with the 6th largest economy by nomial GDP, as of October 2024 [3], so it is reasonable to assume that most of these homes are built from bricks.

We assume that the heatwave is the only natural disaster occuring throughout the entire 24 hours.

Justification: If another natural or manmade disaster happens to occur during that 24 hour time frame the temperatures could fluctuate greatly without pattern, indeed if a heatwave occurs at the same time as another natural disaster it would introduce too large of an uncertainty in the accuracy of our model, Hence we simplify our task by assuming that the only disaster is the heatwave.

We assume that the rate of global warming is constant.

Justification: Since the rate of global warming is dependent on numerous factors that are outside of our control we choose to assume that the rate of global warming is a constant value.

2.3 Analysing the problem

Indeed the temperature inside a given home will vary based on a number of environmental and non-environmental factors. It is well known that temperatures throughout the day naturally peak around midday and tend to fall off in the early morning and late at night. This motivates an intuition for a model with a shape similar to the bell curve shape of a normal distribution.

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{1}{2}\left(\frac{x-\mu}{\sigma}\right)^2\right)$$

However the behaviour of such a curve is not so easy to modify when we are given a good number of variables, like in this problem. We instead make use of *multivariate polynomial regression* to find a polynomial that can accurately model the variation of temperatures in a room given n parameters, using a Python program.

2.4 Variables and parameters

Table 1: Variables and parameters.

Name	Notation	Description
Outside temperature	T	Average outside temperature measured in degrees celcius.
Dew point	D	The outside dew point measured in degrees celcius.
Humidity	ζ	Relative humidity expressed as a percentage, where 100% means the air is fully saturated with moisture, and values at 0% indicate the air is not saturated. We let 100% be 1 and 0% be 0 such that $\zeta \in [0, 1]$.
Time	t	The time in seconds after 00:00 (midnight).
Year	Y	The current year.

2.5 Developing the model

Define $\theta(T, D, \zeta, t)$ to be the function that predicts an indoor temperature during a heatwave in Birmingham, England. Since we have 4 variables it is easy to see that our function will have $\binom{n+4-1}{4-1}$ terms as a consequence of stars and bars. So θ is a multivarite polynomial of degree n the following form,

$$\sum_{\substack{e_T + e_D + \ldots + e_t = n \\ 1 \le r \le \binom{n+4}{4}}} \lambda_r T^{e_T} D^{e_D} \cdots t^{e_t}$$

Indeed global warming is causing the average temperature to increase each year. According to the *National Oceanic and Atmospheric Administration* (NOAA) earth's average temperature has been increasing by 0.06 degrees celcius each decade since 1850 [4]. So we make use of our assumption that the rate of global warming is constant and redefine θ to be that summation plus $(Y - 2025)0.006$.

A first degree polynomial would be to simple and hence unsuitable to model the curves and other features that we might find in a graph of temperature against time. A second degree polynomial may be able to capture basic trends in the data. Using a polynomial of a very high degree will cause the $\binom{n+3}{3}$ term to grow extremely quickly, meaning that an algorithm would to find our coefficients λ_i would run very slowly. Furthermore usage of a high polynomial degree could potentially lead to our model *overfitting* to our dataset, potentially reducing the final accuracy of our model. So instead we simply use a degree 3 polynomial. We fit our model against the data given in the data set provided by M3. [2]

We use linear regression to find our coeffecients with the Python Scikit-learn library, and then plot the residuals against the data set that we trained against to evaluate the performance of our model. Here is a snippet of the code. *See the appendix for the entirety of the code.*

```
1  # Degree of our polynomial
2  DEG = 3
```

```python
3
4    # Data given by M3
5    data = {
6        ...
7    }
8
9    df = pd.DataFrame(data)
10
11   ...
12
13   # Create a pipeline that first transforms the features into polynomial features
14   # then fits a linear regression model
15   model = make_pipeline(PolynomialFeatures(degree=DEG, include_bias=False),
     ↪  LinearRegression())
16
17   # Fit the model
18   model.fit(X, y)
19
20   # To see the learned coefficients, we can extract them from the LinearRegression
     ↪   step
21   lin_reg = model.named_steps['linearregression']
22   print("Intercept:", lin_reg.intercept_)
23   print("Coefficients:", lin_reg.coef_)
24
25   # Predict on the entire dataset to test it
26   y_pred_all = model.predict(X)
27   residuals = y - y_pred_all
28
29   ...
30   plt.show()
```

Here are the outputs, the predicted temperatures start from 00:00 and increase by 1 hour each time. So the first value gives temperature in degrees celcius at 00:00 then the second gives temperature at 01:00 and so on.

```
Intercept: -1275787.4333442824
Coefficients: [-6.23973942e-23 -1.96969079e-19 -6.91442636e-19  2.46451279e-13
  4.92628196e-17  3.53635378e-18 -6.75490063e-21 -3.27335745e-22
  9.75228429e-26  4.16915770e-23  1.93397873e-19 -2.40666684e-18
  2.66786448e-11  1.29282060e-12 -3.88232972e-16 -2.64492896e-19
 -1.08428510e-20 -2.27176515e-22  5.27294229e-24]

Predicted Temperatures: [21.76129363 18.76824161 17.05217446 17.06800382 16.46977547 16.485595
  17.18668672 27.66137863 29.69367947 33.46963231 34.61344218 35.62198628
  35.65301066 35.78871563 35.5220175  34.25571878 34.2714295  33.68865326
```

```
32.57605841 32.59175584 27.90601002 27.9216973   27.93737816 26.53566475]
```
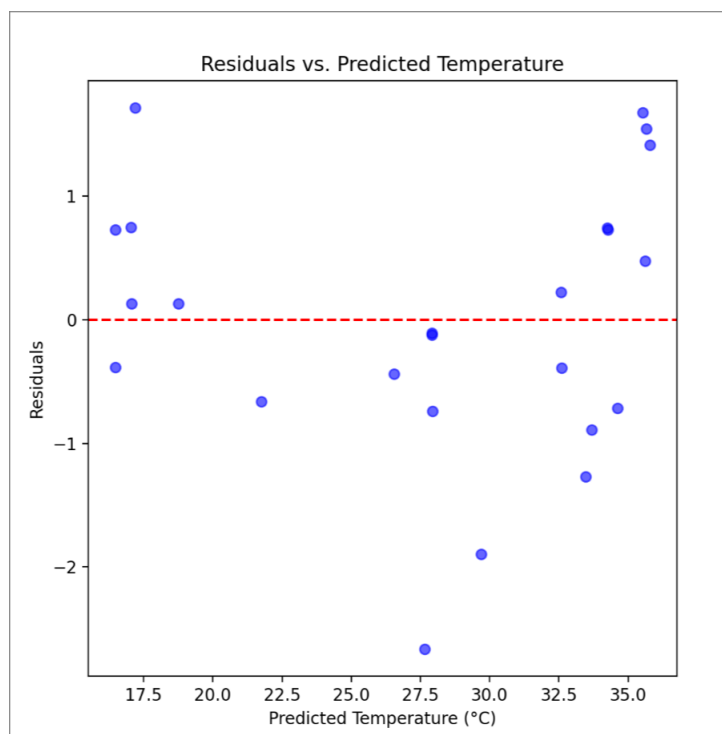

Figure 1: Plot of the residuals against the data set given.

2.6 Discussion of results

By analysing our Matplotlib graph, we find that we have a discrepency of at most 1-2 degrees celcius when predicting temperatures. This demonstrates that our model produces relatively accurate values with minimal error. The strengths of this model include the fact that it can potentially be improved significantly by simply changing the degree that we use in our approximation, and that we are able to produce a clear visualisation to present the accuracy of our predicitons each time. Since our output coefficients ended up being rather small (most of them were of the order 10^{-11} or smaller) this could introduce a measure of numerical instability in our model. Since computers work with a limited amount of precision it would certainly be preferable to have coefficients that require less precision in order to be represented.

3 Q2: Power hungry

3.1 Problem statement

For Q2 we interpret the problem statement like so:

Predict the peak demand of power in kWh that the city's power grid should be prepared to handle during the summer months, now and in 20 years from now.

3.2 Assumptions

We assume that no disaster occurs in the time period that we are predicting.

Justification: Disasters both manmade and natural will cause massive fluctuations in the amount of power used, whilst we have no information to predict when future disasters will occur. For this reason it makes sense to make such an assumption.

We assume that the proportion of electricity consumed in summer months per an entire year is a constant.

Justification: Since we only have data for each year but we instead would like data for the summer months. Since we don't have data for other months we have to assume that the proportion of energy consumption per year used in the summer months is a constant value, χ.

3.3 Analysing the problem

By studying the dataset given by M3 [2], we find that the domestic consumption of electricity seems to be decreasing at a constant rate with minor fluctuations in between. The non-domestic consumption of electricity seems to have remained relatively constant from 2012 to 2022 this time with larger fluctuations from 2015 to 2019.

Both of these data sets demonstrate an amount of fluctuation in their overall trends. This motivates an attempt at using a model that places more emphasis on recent data rather than using a model given by an explicit well defined mathematical function that would place more emphasis on the long term trends.

In order to achive this we make use of the *Holt-Winters Exponential* smoothing model with an additive time series. This model takes into account both long term progression (trends) and short term patterns and cycles (seasonality) as well as the baseline value of our model (level). The model would be of the following form, for δ time steps in the future the energy used will be given by

$$E_{t+\delta} = l_t + \delta b_t + s_{t-T+\delta}$$

where l_t is the level at time t, b_t is the trend at time t, s_t is the season at time t, E_t is the energy consumption at time t and T is the seasonal period.

We create 2 models one to model the variation in the domestic energy consumption, and one to model the variations in the non-domestic energy consumption, then sum up these models in order to give a value for resultant energy consumption.

3.4 Variables and parameters

Table 2: Variables and parameters.

Name	Notation	Description
Constant of proportionality	χ	The constant of proportionality of the consumed energy in summer per the total energy consumed per year.
Level, Domestic energy	$l_{d,t}$	The level of the domestic energy at a time t.
Level, Industrial energy	$l_{i,t}$	The level of the industrial energy at a time t.
Trend, Domestic energy	$b_{d,t}$	The trend of the domestic energy at a time t.
Trend, Industrial energy	$b_{i,t}$	The trend of the industrial energy at a time t.
Season, Domestic energy	$s_{d,t}$	The season of the domestic energy at a time t.
Season, Industrial energy	$s_{i,t}$	The season of the industrial energy at a time t.
Year	Y	The year of when we want to predict electricity consumed

3.5 Developing the model

Let $E_d(t)$ be a function of domestic energy consumption for the year t in kWh. Let $E_i(t)$ be a function of industrial (non-domestic) energy consumption for the year t in kWh. We define $\Sigma(t)$ to be the function that gives the total energy used at a time t, such that $\Sigma(t) = \chi(E_d(t) + E_i(t))$. We make use of the data set given by the M3 [2], which only gives energy consumption per year in Birmingham, England except for 2019 where energy consumption is broken down by month.

By computing the proportion of energy consumed in the summer months of 2019 against the total energy used in 2019 we determine a value for $\chi \approx 0.081$ that is 8.1%.

We use the statsmodels library in Python to implement this algorithm, here is a snippet of the algorithm: *Full code can be viewed in the appendix.*

```
1   ...
2   # Constant of propotionality (refer to paper)
3   CHI = 0.081
4   data["Domestic Consumption"]     *= CHI
5   data["Non-Domestic Consumption"] *= CHI
6
7   ...
8   # Fit Holt-Winters models for each consumption type
9   domestic_model = ExponentialSmoothing(
10      data["Domestic Consumption"],
11      trend='add',      # Additive trend since the data is annual
12      seasonal=None,    # No seasonal component for annual data
13      initialization_method="estimated"
14  ).fit()
15
```

```
16  nondomestic_model = ExponentialSmoothing(
17      data["Non-Domestic Consumption"],
18      trend='add',
19      seasonal=None,
20      initialization_method="estimated"
21  ).fit()
22
23  # Compute the fitted values for each model and sum them to get the combined
    ↪  (total) fitted values
24  data["Total Fitted"] = domestic_model.fittedvalues +
    ↪  nondomestic_model.fittedvalues
25
26  # Also, calculate the actual total consumption for comparison
27  data["Total Actual"] = data["Domestic Consumption"] + data["Non-Domestic
    ↪  Consumption"]
28
29  # Forecast the next 20 years using both models and sum their forecasts
30  forecast_steps = 20
31  domestic_forecast = domestic_model.forecast(steps=forecast_steps)
32  nondomestic_forecast = nondomestic_model.forecast(steps=forecast_steps)
33  total_forecast = domestic_forecast + nondomestic_forecast
34
35  ...
36
37  plt.show()
```

Here is the output from the code:

```
Forecast for Total Consumption for the next 20 years:
1     2.765753e+08
2     2.761579e+08
3     2.757406e+08
4     2.753233e+08
5     2.749060e+08
6     2.744886e+08
7     2.740713e+08
8     2.736540e+08
9     2.732366e+08
10    2.728193e+08
11    2.724020e+08
12    2.719846e+08
13    2.715673e+08
14    2.711500e+08
15    2.707327e+08
16    2.703153e+08
```

```
17     2.698980e+08
18     2.694807e+08
19     2.690633e+08
20     2.686460e+08
```
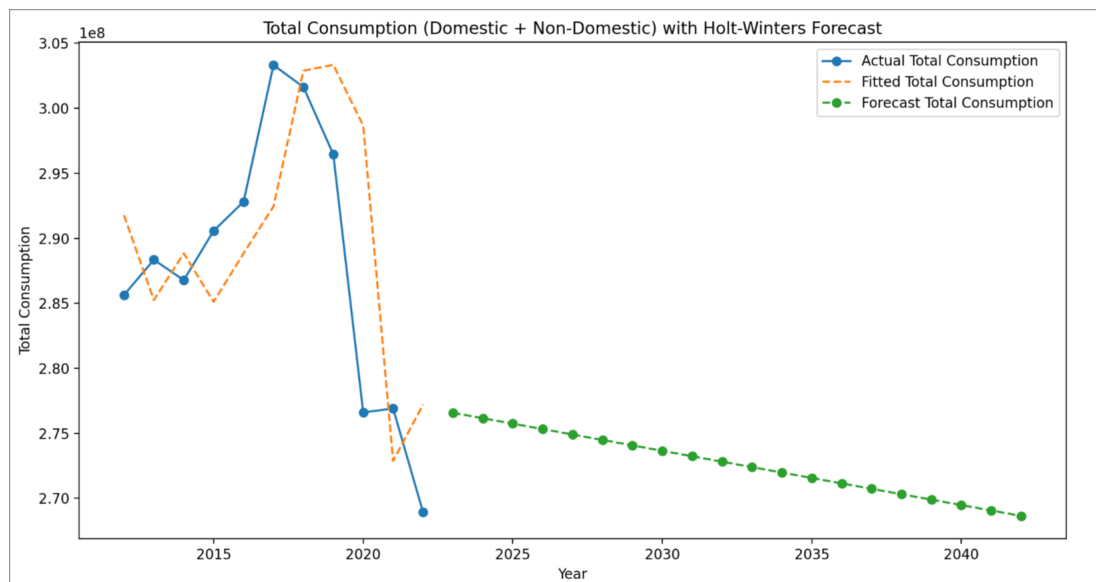

Figure 2: Plot of the actual total energy consumption with our modelled total energy consumption and forecast total energy consumption in summer months.

3.6 Discussion of results

From analysing the plots of the graphs in figure 2, we can see that our fitted model is rather accurately predicting the total energy consumption in summer months from 2012 to 2022, with a small amount of lag. After 2022 our forecast prediction continues the decreasing trend demonstrated in the last few years, without any of the fluctuations shown in the data set.

This model accurately captured the general direction of the data set. However since the forecasted data doesn't have the same fluctuations as the data set given, it doesn't seem unreasonable to conclude that the accuracy and certainty of our predictions could be considered dubious.

4 Q3: Beat the Heat

4.1 Problem statement

In Q3 we interpret the problem statement like so:

Develop a system to determine a vulnerability score which can be assigned to a variety of neighborhoods. Depending on the score, the neighborhoods will be allocated a sufficient amount of resources to reduce the effects of a heat wave or power grid failure. Include an approach of how Birmingham can use these scores in their process of managing heat waves.

4.2 Assumptions

We assume that the wind speed is constant

Justification: Since the rate at which a person cools is directly proportional to the wind speed, i.e. the faster the wind speed the faster a person will be able to cool. Thus, in order to eradicate this parameter and simplify the problem this will be assumed to be constant.

We assume that the people in Birmingham are static, that is they have negligable movement throughout the entire 24 hours

Justification: A person's body temperature will heavily affect the heat index which will be calculated so by minimizing ones movement the value of the heat index will be of higher accuracy, in other words it will provide a better representation of how hot a particular person feels throughout the 24 hour heat wave in Birmingham.

We assume that the power outages will occur during the summer months

Justification: By applying this assumption within our model we will be able to effectively use the Heat Index as a vulnerability parameter.

4.3 Analysing the problem

By studying the following article, *'Predicting Indoor Heat Exposure Risk during Extreme Heat Events'* [1], it was clear that we would use the Heat Index as a suitable vulnerability score. The Heat Index effectively quantifies how hot a human body will feel at a particular temperature and humidity [1]. Thus, it is an important factor in terms of determining what areas will need additional support and resources during a heat wave. It is also an effective parameter to minimize effects during power outages. For example, by observing a variety of Heat Index values at particular points in time during a day we will be able to determine what parts of the city are being affected by hot temperature the most and allocate a suitable amount of power to this part of the city rather than other parts during these times and at other times other parts of the city may be affected more - so suitable resources can be allocated to these parts instead.

In order to achieve this we will make use of the Heat Index Equation where the key assumptions for this equation have been explained above. The equation is as follows:

$$\begin{aligned}
\gamma = {}& -42.379 + 2.04901523\,T + 10.14333127\,\text{RH} \\
& - 0.22475541\,T\,\text{RH} - 0.00683783\,T^2 \\
& - 0.05481717\,\text{RH}^2 + 0.00122874\,T^2\,\text{RH} \\
& + 0.00085282\,T\,\text{RH}^2 - 0.00000199\,T^2\text{RH}^2
\end{aligned}$$

However, we need to calculate the Relative Humidity for each value of Humidity given, this is calculated by the following:

$$\epsilon = 100 \times \left[\frac{\exp\left(\frac{17.625 D_p}{243.04 + D_p}\right)}{\exp\left(\frac{17.625 T}{243.64 + T}\right)} \right]$$

4.4 Variables and parameters

Table 3: Variables and parameters.

Name	Notation	Description
Relative Humidity	ϵ	Measure of the water content in the air compared to the maximum water content that the air can hold at a particular temperature.
Temperature	T	Temperature of the air in degrees celcius.
Heat index	γ	A measure of how hot the air temperature feels at a particular temperature and humidity.
Dew point temperature	ρ	The temperature in degrees celcius at which air holds the maximum amount of water vapour

4.5 Developing the model

Define $\gamma(\epsilon, T)$ to be the function that calculates the heat index at any temperature, T, and relative humidity, ϵ. Let $\epsilon(\rho, T)$ to be the function that calculates the relative humidity at a particular value of ρ and T. Again, we will make use of the data set given in the M3 [2], which gives the values for ρ and T over the 24 hour time period of the heat wave.

Here is a snippet of the algorithm: *Full code can be viewed in appendix*

```python
import math
import matplotlib.pyplot as plt

# Calculating Relative Humidity
def CalcHumidity(airTemp, dewPointTemp):
    top = math.exp( (17.625 * dewPointTemp) / (243.04 + dewPointTemp) )
    bottom = math.exp( (17.625 * airTemp) / (243.04 + airTemp) )
    newHumidity = (top / bottom) * 100
    return newHumidity

# Calculating corresponding Heat Index
def CalcHeatIndex(airTemp, newHumidity):
    calculatedIndex = -42.379 + 2.04901523 * airTemp + 10.14333127 * newHumidity
        - 0.22475541 * airTemp * newHumidity - 0.00683783 * (airTemp**2) -
        0.05481717 * (newHumidity**2) + 0.00122874*(airTemp**2) * newHumidity +
        0.00085282 * airTemp * (newHumidity**2) - 0.00000199 * (airTemp **2) *
        (newHumidity**2)
```

```
14        return calculatedIndex
15
16    ...
17    plt.show()
```

Here is the output of the code:

Figure 3: Plot of the heat index against time.

5 Appendix

5.1 Q1

```
1    import pandas as pd
2    from sklearn.preprocessing import PolynomialFeatures
3    from sklearn.linear_model import LinearRegression
4    from sklearn.pipeline import make_pipeline
5    import matplotlib.pyplot as plt
6
7    # Degree of our polynomial
8    DEG = 3
9
10   # Data given by M3
```

```
11   data = {
12       "Time": ["12:00 AM", "1:00 AM", "2:00 AM", "3:00 AM", "4:00 AM", "5:00 AM",
         ↪  "6:00 AM", "7:00 AM", "8:00 AM", "9:00 AM",
13               "10:00 AM", "11:00 AM", "12:00 PM", "1:00 PM", "2:00 PM", "3:00 PM",
                 ↪  "4:00 PM", "5:00 PM", "6:00 PM", "7:00 PM",
14               "8:00 PM", "9:00 PM", "10:00 PM", "11:00 PM"],
15       "Temperature (°C)": [21.1, 18.9, 17.8, 17.2, 17.2, 16.1, 18.9, 25.0, 27.8,
         ↪  32.2, 33.9, 36.1, 37.2, 37.2, 37.2, 35.0,
16                           35.0, 32.8, 32.8, 32.2, 27.8, 27.8, 27.2, 26.1],
17       "Dew Point (°C)": [12.8, 12.8, 12.2, 12.2, 11.1, 11.1, 13.9, 12.2, 11.1,
         ↪  11.1, 11.1, 12.2, 10.0, 11.1, 12.8, 13.9, 13.9,
18                         12.8, 12.8, 12.8, 15.0, 15.0, 15.0, 16.1],
19       "Humidity (%)": [60, 68, 72, 72, 72, 72, 73, 44, 37, 27, 24, 23, 19, 21, 24,
         ↪  28, 28, 29, 32, 32, 45, 45, 45, 48]
20   }
21
22   df = pd.DataFrame(data)
23
24   # We want times to be seconds after 00:00
25   df["Time"] = pd.to_datetime(df["Time"], format="%I:%M %p")
26   df["Time"] = (df["Time"] - pd.Timestamp("00:00:00")).dt.total_seconds()
27   df["Time"] = df["Time"].astype(int)
28
29   # Humidity % values should be in the interval [0, 1]
30   df["Humidity (%)"] = df["Humidity (%)"] / 100
31
32   # Prepare features (X) and target (y)
33   # Here, we predict Temperature based on Time, Dew Point, and Humidity
34   X = df[["Time", "Dew Point (°C)", "Humidity (%)"]]
35   y = df["Temperature (°C)"]
36
37   # Create a pipeline that first transforms the features into polynomial features
38   # then fits a linear regression model
39   model = make_pipeline(PolynomialFeatures(degree=DEG, include_bias=False),
     ↪  LinearRegression())
40
41   # Fit the model
42   model.fit(X, y)
43
44   # To see the learned coefficients, we can extract them from the LinearRegression
     ↪  step
45   lin_reg = model.named_steps['linearregression']
46   print("Intercept:", lin_reg.intercept_)
47   print("Coefficients:", lin_reg.coef_)
48
```

```python
49  def Predict_Temps(new_data, year):
50      # Predict on the given new_data
51      y_pred_all = model.predict(new_data)
52
53      # Apply the year adjustment
54      y_pred_all_adjusted = y_pred_all + (year - 2025) * (0.006)
55
56      # Print the predicted temperatures
57      print("Predicted Temperatures:", y_pred_all_adjusted)
58
59      # Calculate residuals using the original y for comparison
60      residuals = y - y_pred_all_adjusted
61
62      plt.figure(figsize=(10, 6))
63      plt.scatter(y_pred_all_adjusted, residuals, color='blue', alpha=0.6)
64      plt.axhline(0, color='red', linestyle='--')
65      plt.xlabel("Predicted Temperature (°C)")
66      plt.ylabel("Residuals")
67      plt.title("Residuals vs. Predicted Temperature")
68      plt.show()
69
70  Predict_Temps(X, 2025)
```

5.2 Q2

```python
1   import pandas as pd
2   import matplotlib.pyplot as plt
3   from statsmodels.tsa.holtwinters import ExponentialSmoothing
4
5   # Create the DataFrame with the provided data
6   data = pd.DataFrame({
7       "Year": [2022, 2021, 2020, 2019, 2018, 2017, 2016, 2015, 2014, 2013, 2012],
8       "Domestic Consumption": [1424117394, 1525486540, 1609229857, 1535046668,
        ↪ 1551441077,
9                               1584877006, 1602218919, 1632646065, 1650323202,
                                ↪ 1632305132, 1663571950],
10      "Non-Domestic Consumption": [1896047335, 1893143269, 1805655202, 2124916982,
        ↪ 2172490874,
11                                  2159883548, 2012774683, 1954766502, 1890155436,
                                    ↪ 1927600374, 1862786430]
12  })
13
14  # Constant of propotionality (refer to paper)
```

```
15  CHI = 0.081
16  data["Domestic Consumption"]     *= CHI
17  data["Non-Domestic Consumption"] *= CHI
18
19  # Sort the data by Year and set Year as index
20  data = data.sort_values("Year")
21  data.set_index("Year", inplace=True)
22
23  # Fit Holt-Winters models for each consumption type
24  domestic_model = ExponentialSmoothing(
25      data["Domestic Consumption"],
26      trend='add',      # Additive trend since the data is annual
27      seasonal=None,    # No seasonal component for annual data
28      initialization_method="estimated"
29  ).fit()
30
31  nondomestic_model = ExponentialSmoothing(
32      data["Non-Domestic Consumption"],
33      trend='add',
34      seasonal=None,
35      initialization_method="estimated"
36  ).fit()
37
38  # Compute the fitted values for each model and sum them to get the combined
    ↪  (total) fitted values
39  data["Total Fitted"] = domestic_model.fittedvalues +
    ↪  nondomestic_model.fittedvalues
40
41  # Also, calculate the actual total consumption for comparison
42  data["Total Actual"] = data["Domestic Consumption"] + data["Non-Domestic
    ↪  Consumption"]
43
44  # Forecast the next 20 years using both models and sum their forecasts
45  forecast_steps = 20
46  domestic_forecast = domestic_model.forecast(steps=forecast_steps)
47  nondomestic_forecast = nondomestic_model.forecast(steps=forecast_steps)
48  total_forecast = domestic_forecast + nondomestic_forecast
49
50  # Create a new index for the forecast years (assuming years continue
    ↪  consecutively)
51  forecast_years = list(range(data.index[-1] + 1, data.index[-1] + forecast_steps
    ↪  + 1))
52
53  # Plot the actual total consumption, the fitted total consumption, and the
    ↪  forecast
```

```
54  plt.figure(figsize=(12, 6))
55  plt.plot(data.index, data["Total Actual"], label="Actual Total Consumption",
    ↪  marker='o')
56  plt.plot(data.index, data["Total Fitted"], label="Fitted Total Consumption",
    ↪  linestyle="--")
57  plt.plot(forecast_years, total_forecast, label="Forecast Total Consumption",
    ↪  marker='o', linestyle="--")
58  plt.xlabel("Year")
59  plt.ylabel("Total Consumption")
60  plt.title("Total Consumption (Domestic + Non-Domestic) with Holt-Winters
    ↪  Forecast")
61  plt.legend()
62  plt.show()
63
64  # Print the forecast values
65  print("Forecast for Total Consumption for the next 20 years:")
66  print(total_forecast)
```

5.3 Q3

```
1   import pandas as pd
2   import matplotlib.pyplot as plt
3   from statsmodels.tsa.holtwinters import ExponentialSmoothing
4
5   # Create the DataFrame with the provided data
6   data = pd.DataFrame({
7       "Year": [2022, 2021, 2020, 2019, 2018, 2017, 2016, 2015, 2014, 2013, 2012],
8       "Domestic Consumption": [1424117394, 1525486540, 1609229857, 1535046668,
        ↪  1551441077,
9                                 1584877006, 1602218919, 1632646065, 1650323202,
                                  ↪  1632305132, 1663571950],
10      "Non-Domestic Consumption": [1896047335, 1893143269, 1805655202, 2124916982,
        ↪  2172490874,
11                                 2159883548, 2012774683, 1954766502, 1890155436,
                                   ↪  1927600374, 1862786430]
12  })
13
14  # Constant of propotionality (refer to paper)
15  CHI = 0.081
16  data["Domestic Consumption"]     *= CHI
17  data["Non-Domestic Consumption"] *= CHI
18
19  # Sort the data by Year and set Year as index
```

```python
20    data = data.sort_values("Year")
21    data.set_index("Year", inplace=True)
22
23    # Fit Holt-Winters models for each consumption type
24    domestic_model = ExponentialSmoothing(
25        data["Domestic Consumption"],
26        trend='add',        # Additive trend since the data is annual
27        seasonal=None,      # No seasonal component for annual data
28        initialization_method="estimated"
29    ).fit()
30
31    nondomestic_model = ExponentialSmoothing(
32        data["Non-Domestic Consumption"],
33        trend='add',
34        seasonal=None,
35        initialization_method="estimated"
36    ).fit()
37
38    # Compute the fitted values for each model and sum them to get the combined
      ↪  (total) fitted values
39    data["Total Fitted"] = domestic_model.fittedvalues +
      ↪  nondomestic_model.fittedvalues
40
41    # Also, calculate the actual total consumption for comparison
42    data["Total Actual"] = data["Domestic Consumption"] + data["Non-Domestic
      ↪  Consumption"]
43
44    # Forecast the next 20 years using both models and sum their forecasts
45    forecast_steps = 20
46    domestic_forecast = domestic_model.forecast(steps=forecast_steps)
47    nondomestic_forecast = nondomestic_model.forecast(steps=forecast_steps)
48    total_forecast = domestic_forecast + nondomestic_forecast
49
50    # Create a new index for the forecast years (assuming years continue
      ↪  consecutively)
51    forecast_years = list(range(data.index[-1] + 1, data.index[-1] + forecast_steps
      ↪  + 1))
52
53    # Plot the actual total consumption, the fitted total consumption, and the
      ↪  forecast
54    plt.figure(figsize=(12, 6))
55    plt.plot(data.index, data["Total Actual"], label="Actual Total Consumption",
      ↪  marker='o')
56    plt.plot(data.index, data["Total Fitted"], label="Fitted Total Consumption",
      ↪  linestyle="--")
```

```python
57  plt.plot(forecast_years, total_forecast, label="Forecast Total Consumption",
    ↪  marker='o', linestyle="--")
58  plt.xlabel("Year")
59  plt.ylabel("Total Consumption")
60  plt.title("Total Consumption (Domestic + Non-Domestic) with Holt-Winters
    ↪  Forecast")
61  plt.legend()
62  plt.show()
63
64  # Print the forecast values
65  print("Forecast for Total Consumption for the next 20 years:")
66  print(total_forecast)
```

References

[1] Department of Environmental Health Sciences, Columbia University, New York. *Predicting Indoor Heat Exposure Risk During Extreme Heat Events*. 2015. URL: https://pmc.ncbi.nlm.nih.gov/articles/PMC4121079/#:~:text=Modeling%20the%20associations%20between%20indoor,thresholds%20during%20these%20extreme%20events..

[2] Hot Button Issue, curated data. *MathWorks Math Modeling Challenge 2025*. 2025. URL: https://m3challenge.siam.org/897bjhb54cgfc/.

[3] International Monetary Fund. *World Economic Outlook*. 2024. URL: https://www.imf.org/external/datamapper/profile/GBR.

[4] National centers for Environmental Information. *Global Climate Report*. 2023. URL: https://www.ncei.noaa.gov/access/monitoring/monthly-report/global/202313.